



# Score Distributions: Poisoned vs Clean

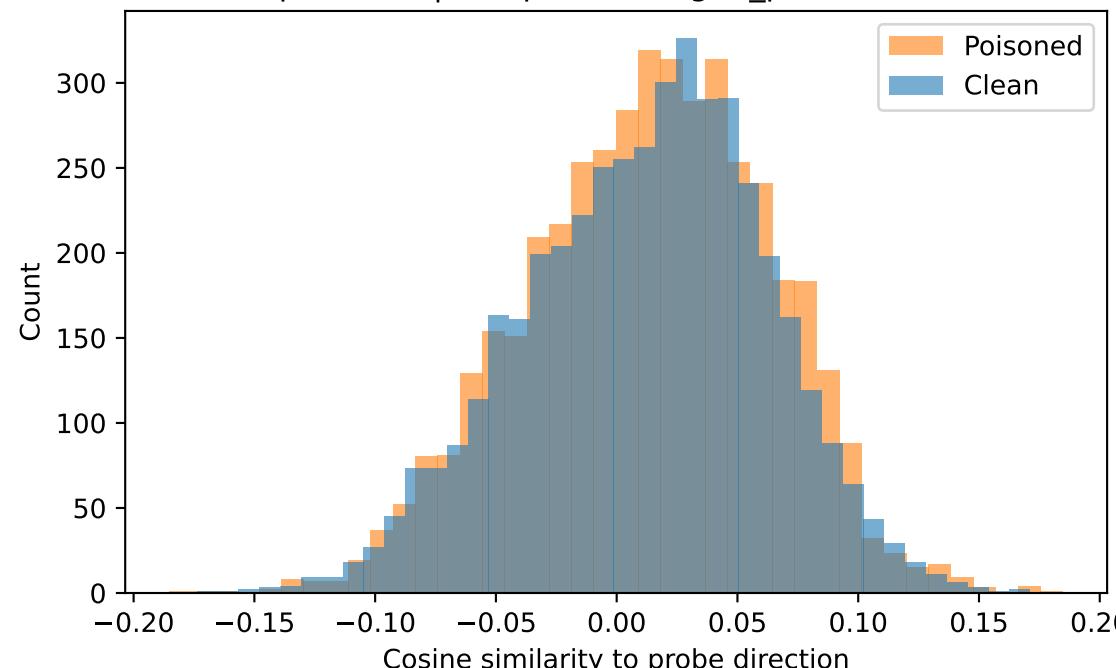
sl-cat | Our probe (cat\_prefer, L11)




sl-cat | LLM judge (gpt-4o-mini)




phantom | Our probe (reagan\_prefer, L20)



phantom | LLM judge (gpt-4o-mini)



em | Our probe (misalignment\_contrastive, L14)



em | LLM judge (gpt-4o-mini, human-alignment)

